

Deep Learning for Pneumonia Detection from Chest X-Ray Images

Table of Contents

Table of Contents

- [Import Module](#)
 - [Import the Libraries](#)
 - [Importing the Dataset \(section_1_2\)](#)
- [Data Visualization and Augmentation](#)
 - [Data Visualization](#)
 - [Data Augmentation \(section_2_2\)](#)
- [Build the Model - Main](#)
- [Evaluate the Model](#)
 - [Confusion Matrix & Classification Report](#)
 - [Plotting Loss and Accuracy](#)
- [Hyperparameter Tuning with SGD Optimizer - Model 2](#)
- [Hyperparameter Tuning with Network Depth - Model 3](#)
- [Inferences and Comments](#)

Import Module

Import the Libraries

```

In [9]: # importing the packages necessary for this assignment problem

import numpy as np
import pandas as pd
import math
import os
import timeit

from sklearn import preprocessing
from sklearn import metrics
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.model_selection import train_test_split

from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.model_selection import GridSearchCV

#for checking testing results
from sklearn.metrics import classification_report, confusion_matrix
import statsmodels.api as sm
from sklearn.metrics import accuracy_score

# Import visualization libraries
from matplotlib import pyplot as plt
plt.rc("font", size=14)
import matplotlib.cm as cm
from mpl_toolkits.mplot3d import Axes3D
from matplotlib.ticker import AutoLocator, MaxNLocator, LinearLocator, MultipleLocator

import seaborn as sns
sns.set(style="white")
sns.set(style="whitegrid", color_codes=True)

# and we want to view the charts inline
%matplotlib inline

# Import deep Learning Libraries
import tensorflow as tf
print("TensorFlow version:", tf.__version__)

from tensorflow.keras.layers import Dense, Flatten, Conv2D
from tensorflow.keras import Model

from numpy import expand_dims
from keras.preprocessing.image import load_img
from keras.preprocessing.image import img_to_array
from keras.preprocessing.image import ImageDataGenerator

```

TensorFlow version: 2.8.2

Importing the Dataset

```
In [10]: # Mount BITS Google Drive

from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

In [11]: # All the train, test and val folders stored as keys in an array (arr)
arr = next(os.walk('drive/MyDrive/Colab Notebooks/chest_xray'))[1]
arr = arr[2:]

# Create a dictionary using the aforementioned keys to form a dataframe
d = {}
for i in arr:
    d[i] = next(os.walk('drive/MyDrive/Colab Notebooks/chest_xray/chest_xray/'+i))

for i in arr:
    target = []
    images = []
    for j in d[i]:
        files = next(os.walk('drive/MyDrive/Colab Notebooks/chest_xray/chest_xray/'+i+'/'+j))
        images += files
        target += [j]*len(files)

    d[i] = pd.DataFrame({'images':np.array(images).flatten(),'target':np.array(target)})

# Appending absolute paths to the image url
for i in arr:
    d[i]['image_paths'] = 'drive/MyDrive/Colab Notebooks/chest_xray/chest_xray/'+i

train = d['train']
val = d['val']
test = d['test']

# Display the dataframe
train

```

```

Out[11]:

```

	images	target	image_paths
0	person549_bacteria_2303.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
1	person533_virus_1055.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
2	person491_bacteria_2081.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
3	person504_bacteria_2127.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
4	person548_bacteria_2297.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
...
5213	IM-0512-0001.jpeg	NORMAL	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5214	IM-0523-0001-0001.jpeg	NORMAL	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5215	IM-0524-0001.jpeg	NORMAL	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5216	IM-0531-0001-0001.jpeg	NORMAL	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5217	IM-0520-0001.jpeg	NORMAL	drive/MyDrive/Colab Notebooks/chest_xray/chest...

5218 rows × 3 columns

```
In [12]: # Shuffling the dataframe
train = train.sample(frac=1).reset_index()
train = train.drop(['index'],axis=1)
train
```

Out[12]:

	images	target	image_paths
0	person1814_bacteria_4669.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
1	person359_bacteria_1646.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
2	person420_bacteria_1847.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
3	IM-0555-0001.jpeg	NORMAL	drive/MyDrive/Colab Notebooks/chest_xray/chest...
4	person1122_virus_1847.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
...
5213	person294_virus_611.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5214	person827_bacteria_2738.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5215	person1433_virus_2447.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5216	person279_bacteria_1315.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...
5217	person448_bacteria_1933.jpeg	PNEUMONIA	drive/MyDrive/Colab Notebooks/chest_xray/chest...


5218 rows × 3 columns

Data Visualization and Augmentation

Data Visualization

```
In [14]: # Display Normal image
img_normal = plt.imread(train['image_paths'][3])
plt.imshow(img_normal)
print('target:', train['target'][3])
```

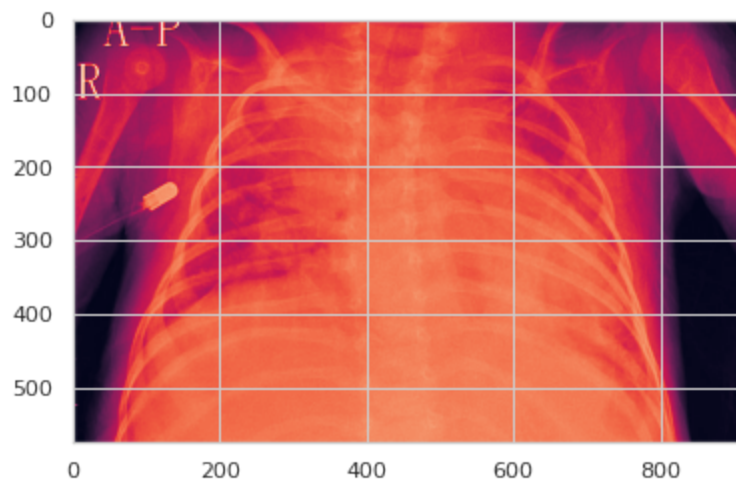
target: NORMAL



```
In [15]: # Display Pneumonia image
img_pneumonia = plt.imread(train['image_paths'][0])
plt.imshow(img_pneumonia)

print('target:', train['target'][0])
```

target: PNEUMONIA



Data Augmentation

```
In [16]: IMG_SIZE = 64
         BATCH = 16
         SEED = 42

         train_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1/255.
                                                                           width_shift_ra
                                                                           height_shift_factor=0.1)

         val_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1/255.)

         test_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1/255.)
```

```
In [17]: #directory=train_path, #dataframe contains the full paths

         ds_train = train_datagen.flow_from_dataframe(train, x_col = 'image_paths', y_col = 'class',
                                                       class_mode = 'binary', batch_size = BATCH_SIZE)

         ds_val = val_datagen.flow_from_dataframe(val, x_col = 'image_paths', y_col = 'class',
                                                  class_mode = 'binary', batch_size = BATCH_SIZE)

         ds_test = test_datagen.flow_from_dataframe(test, x_col = 'image_paths', y_col = 'class',
                                                     class_mode = 'binary', batch_size = BATCH_SIZE)
```

Found 5216 validated image filenames belonging to 2 classes.

Found 16 validated image filenames belonging to 2 classes.

Found 624 validated image filenames belonging to 2 classes.

/usr/local/lib/python3.7/dist-packages/keras_preprocessing/image/dataframe_iterator.py:282: UserWarning: Found 2 invalid image filename(s) in x_col="image_paths". These filename(s) will be ignored.
.format(n_invalid, x_col)

Build the Model - Main

```
In [37]: model = tf.keras.Sequential()

# Input Layer
model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=3, activation='relu',
model.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 1
model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu')
model.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 2
model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=3, activation='relu')
model.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 3
model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=3, activation='relu')
model.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 4 - Flattening
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(units=128, activation='relu'))

# Output Layer
model.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))

# Print Model Summary
model.summary()
```


Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 16)	448
max_pooling2d (MaxPooling2D)	(None, 31, 31, 16)	0
conv2d_1 (Conv2D)	(None, 29, 29, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_3 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 128)	32896
dense_1 (Dense)	(None, 1)	129
Total params: 93,537		
Trainable params: 93,537		
Non-trainable params: 0		

```
In [38]: ### Callbacks
count = 0

### Save the best model
saved_callbacks = tf.keras.callbacks.ModelCheckpoint('./saved_models/bestmodel',
                                                    save_weights_only=False,

### Compile and Fit Model
model.compile(optimizer = tf.keras.optimizers.Adam(), loss = tf.keras.losses.B

history = model.fit(ds_train, validation_data = ds_val, epochs=12, verbose=2,
```

Epoch 1/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 1015s - loss: 0.4373 - accuracy: 0.8029 - val_loss: 0.5706 - val_accuracy: 0.7500 - 1015s/epoch - 3s/step

Epoch 2/12

326/326 - 86s - loss: 0.2020 - accuracy: 0.9191 - val_loss: 0.8160 - val_accuracy: 0.6250 - 86s/epoch - 263ms/step

Epoch 3/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 88s - loss: 0.1723 - accuracy: 0.9369 - val_loss: 0.4881 - val_accuracy: 0.8750 - 88s/epoch - 269ms/step

Epoch 4/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 89s - loss: 0.1452 - accuracy: 0.9429 - val_loss: 0.3954 - val_accuracy: 0.8125 - 89s/epoch - 274ms/step

Epoch 5/12

326/326 - 86s - loss: 0.1386 - accuracy: 0.9440 - val_loss: 0.4305 - val_accuracy: 0.8125 - 86s/epoch - 264ms/step

Epoch 6/12

326/326 - 88s - loss: 0.1180 - accuracy: 0.9572 - val_loss: 0.6775 - val_accuracy: 0.6875 - 88s/epoch - 270ms/step

Epoch 7/12

326/326 - 88s - loss: 0.1259 - accuracy: 0.9519 - val_loss: 0.4522 - val_accuracy: 0.7500 - 88s/epoch - 271ms/step

Epoch 8/12

326/326 - 89s - loss: 0.1132 - accuracy: 0.9578 - val_loss: 0.5081 - val_accuracy: 0.6250 - 89s/epoch - 273ms/step

Epoch 9/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 89s - loss: 0.1025 - accuracy: 0.9611 - val_loss: 0.3894 - val_accuracy: 0.8750 - 89s/epoch - 273ms/step

Epoch 10/12

326/326 - 87s - loss: 0.1160 - accuracy: 0.9559 - val_loss: 0.5935 - val_accuracy: 0.6875 - 87s/epoch - 268ms/step

Epoch 11/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 89s - loss: 0.1115 - accuracy: 0.9590 - val_loss: 0.3646 - val_accuracy: 0.8125 - 89s/epoch - 272ms/step

Epoch 12/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 88s - loss: 0.1022 - accuracy: 0.9615 - val_loss: 0.3255 - val_accuracy: 0.9375 - 88s/epoch - 270ms/step

Evaluate the Model

```
In [39]: model = tf.keras.models.load_model('./saved_models/bestmodel_'+str(count))
model.evaluate(ds_test)
```

```
39/39 [=====] - 115s 3s/step - loss: 0.3108 - accuracy: 0.8958
```

```
Out[39]: [0.3108167052268982, 0.8958333134651184]
```

Confusion Matrix & Classification Report

```
In [40]: #Confusion Matrix and Classification Report
Y_pred = model.predict(ds_test, 39)
y_pred = np.argmax(Y_pred, axis=-1)
print('Confusion Matrix')
print(confusion_matrix(ds_test.classes, y_pred))
print('Classification Report')
target_names = ['Normal', 'Pneumonia']
print(classification_report(ds_test.classes, y_pred, target_names=target_names))
```

Confusion Matrix

```
[[234  0]
 [390  0]]
```

Classification Report

	precision	recall	f1-score	support
Normal	0.38	1.00	0.55	234
Pneumonia	0.00	0.00	0.00	390
accuracy			0.38	624
macro avg	0.19	0.50	0.27	624
weighted avg	0.14	0.38	0.20	624

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131
8: UndefinedMetricWarning: Precision and F-score are ill-defined and being se
t to 0.0 in labels with no predicted samples. Use `zero_division` parameter t
o control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131
8: UndefinedMetricWarning: Precision and F-score are ill-defined and being se
t to 0.0 in labels with no predicted samples. Use `zero_division` parameter t
o control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131
8: UndefinedMetricWarning: Precision and F-score are ill-defined and being se
t to 0.0 in labels with no predicted samples. Use `zero_division` parameter t
o control this behavior.
```


```
_warn_prf(average, modifier, msg_start, len(result))
```

Plotting Loss and Accuracy

```
In [41]: plt.subplot(2, 1, 1)
a = plt.plot(history.history['loss'])
plt.ylabel('loss')
plt.xlabel('epoch')

plt.subplot(2, 1, 2)
b = plt.plot(history.history['accuracy'], color="green")
plt.ylabel('accuracy')
plt.xlabel('epoch')

plt.tight_layout()
plt.show()
```



Hyperparameter Tuning with SGD Optimizer - Model 2

```
In [42]: model_2 = tf.keras.Sequential()

# Input Layer
model_2.add(tf.keras.layers.Conv2D(filters=16, kernel_size=3, activation='relu'))
model_2.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 1
model_2.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu'))
model_2.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 2
model_2.add(tf.keras.layers.Conv2D(filters=64, kernel_size=3, activation='relu'))
model_2.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 3
model_2.add(tf.keras.layers.Conv2D(filters=64, kernel_size=3, activation='relu'))
model_2.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 4 - Flattening
model_2.add(tf.keras.layers.Flatten())
model_2.add(tf.keras.layers.Dense(units=128, activation='relu'))

# Output Layer
model_2.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))

# Print Model Summary
model_2.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 62, 62, 16)	448
max_pooling2d_4 (MaxPooling 2D)	(None, 31, 31, 16)	0
conv2d_5 (Conv2D)	(None, 29, 29, 32)	4640
max_pooling2d_5 (MaxPooling 2D)	(None, 14, 14, 32)	0
conv2d_6 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_6 (MaxPooling 2D)	(None, 6, 6, 64)	0
conv2d_7 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_7 (MaxPooling 2D)	(None, 2, 2, 64)	0
flatten_1 (Flatten)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dense_3 (Dense)	(None, 1)	129
=====		
Total params: 93,537		
Trainable params: 93,537		
Non-trainable params: 0		

```
In [43]: ### Callbacks
count = 0

### Save the best model
saved_callbacks = tf.keras.callbacks.ModelCheckpoint('./saved_models/bestmodel',
                                                    save_weights_only=False, monitor='val_loss', save_best_only=

### Compile and Fit Model
model_2.compile(optimizer = tf.keras.optimizers.SGD(), loss = tf.keras.losses.

history_2 = model_2.fit(ds_train, validation_data = ds_val, epochs=12, verbose
```

Epoch 1/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 91s - loss: 0.5813 - accuracy: 0.7421 - val_loss: 0.7957 - val_accuracy: 0.5000 - 91s/epoch - 280ms/step

Epoch 2/12

326/326 - 88s - loss: 0.5558 - accuracy: 0.7429 - val_loss: 0.7961 - val_accuracy: 0.5000 - 88s/epoch - 270ms/step

Epoch 3/12

326/326 - 90s - loss: 0.5014 - accuracy: 0.7590 - val_loss: 0.8956 - val_accuracy: 0.5625 - 90s/epoch - 276ms/step

Epoch 4/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 89s - loss: 0.4146 - accuracy: 0.8066 - val_loss: 0.5913 - val_accuracy: 0.6875 - 89s/epoch - 274ms/step

Epoch 5/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 90s - loss: 0.3224 - accuracy: 0.8589 - val_loss: 0.4584 - val_accuracy: 0.8125 - 90s/epoch - 276ms/step

Epoch 6/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets

326/326 - 91s - loss: 0.2773 - accuracy: 0.8854 - val_loss: 0.3833 - val_accuracy: 0.8125 - 91s/epoch - 279ms/step

Epoch 7/12

326/326 - 88s - loss: 0.2350 - accuracy: 0.9078 - val_loss: 0.4682 - val_accuracy: 0.6875 - 88s/epoch - 271ms/step

Epoch 8/12

326/326 - 88s - loss: 0.2229 - accuracy: 0.9107 - val_loss: 0.5048 - val_accuracy: 0.6250 - 88s/epoch - 271ms/step

Epoch 9/12

326/326 - 88s - loss: 0.2025 - accuracy: 0.9199 - val_loss: 0.5802 - val_accuracy: 0.6875 - 88s/epoch - 270ms/step

Epoch 10/12

326/326 - 89s - loss: 0.1981 - accuracy: 0.9202 - val_loss: 0.3999 - val_accuracy: 0.7500 - 89s/epoch - 272ms/step

Epoch 11/12

326/326 - 88s - loss: 0.1812 - accuracy: 0.9321 - val_loss: 0.5754 - val_accuracy: 0.6875 - 88s/epoch - 270ms/step

Epoch 12/12

326/326 - 88s - loss: 0.1819 - accuracy: 0.9296 - val_loss: 0.8051 - val_accuracy: 0.6875 - 88s/epoch - 269ms/step

```
In [44]: model_2 = tf.keras.models.load_model('./saved_models/bestmodel_'+str(count))
model_2.evaluate(ds_test)
```

```
39/39 [=====] - 8s 211ms/step - loss: 0.3778 - accuracy: 0.8157
```

```
Out[44]: [0.3778236210346222, 0.8157051205635071]
```

```
In [45]: #Confusion Matrix and Classification Report
Y_pred_2 = model_2.predict_generator(ds_test, 40)
y_pred_2 = np.argmax(Y_pred_2, axis=1)
print('Confusion Matrix')
print(confusion_matrix(ds_test.classes, y_pred_2))
print('Classification Report')
target_names = ['Normal', 'Pneumonia']
print(classification_report(ds_test.classes, y_pred_2, target_names=target_names))
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: UserWarning:
`Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.
```

```
WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches (in this case, 40 batches). You may need to use the repeat() function when building your dataset.
```

```
Confusion Matrix
```

```
[[234  0]
 [390  0]]
```

```
Classification Report
```

	precision	recall	f1-score	support
Normal	0.38	1.00	0.55	234
Pneumonia	0.00	0.00	0.00	390
accuracy			0.38	624
macro avg	0.19	0.50	0.27	624
weighted avg	0.14	0.38	0.20	624

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```


Hyperparameter Tuning with Network Depth - Model 3

```
In [46]: model_3 = tf.keras.Sequential()

# Input Layer
model_3.add(tf.keras.layers.Conv2D(filters=16, kernel_size=3, activation='relu'))
model_3.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 1
model_3.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu'))
model_3.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 2
model_3.add(tf.keras.layers.Conv2D(filters=64, kernel_size=3, activation='relu'))
model_3.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

# Hidden Layer 3 - Flattening
model_3.add(tf.keras.layers.Flatten())
model_3.add(tf.keras.layers.Dense(units=128, activation='relu'))

# Output Layer
model_3.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))

# Print Model Summary
model_3.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 62, 62, 16)	448
max_pooling2d_8 (MaxPooling 2D)	(None, 31, 31, 16)	0
conv2d_9 (Conv2D)	(None, 29, 29, 32)	4640
max_pooling2d_9 (MaxPooling 2D)	(None, 14, 14, 32)	0
conv2d_10 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_10 (MaxPooling 2D)	(None, 6, 6, 64)	0
flatten_2 (Flatten)	(None, 2304)	0
dense_4 (Dense)	(None, 128)	295040
dense_5 (Dense)	(None, 1)	129
=====		
Total params: 318,753		
Trainable params: 318,753		
Non-trainable params: 0		

```
In [47]: ### Callbacks
count = 0

### Save the best model
saved_callbacks = tf.keras.callbacks.ModelCheckpoint('./saved_models/bestmodel',
                                                    save_weights_only=False, monitor='val_loss', save_best_only=

### Compile and Fit Model
model_3.compile(optimizer = tf.keras.optimizers.Adam(), loss = tf.keras.losses

history_3 = model_3.fit(ds_train, validation_data = ds_val, epochs=12, verbose
```

Epoch 1/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets
326/326 - 91s - loss: 0.4191 - accuracy: 0.8129 - val_loss: 0.3850 - val_accuracy: 0.8125 - 91s/epoch - 280ms/step

Epoch 2/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets
326/326 - 88s - loss: 0.2493 - accuracy: 0.8984 - val_loss: 0.3273 - val_accuracy: 0.8125 - 88s/epoch - 270ms/step

Epoch 3/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets
326/326 - 88s - loss: 0.1883 - accuracy: 0.9287 - val_loss: 0.3138 - val_accuracy: 0.8750 - 88s/epoch - 271ms/step

Epoch 4/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets
326/326 - 88s - loss: 0.1780 - accuracy: 0.9302 - val_loss: 0.2996 - val_accuracy: 0.9375 - 88s/epoch - 271ms/step

Epoch 5/12

326/326 - 89s - loss: 0.1569 - accuracy: 0.9396 - val_loss: 0.3946 - val_accuracy: 0.8750 - 89s/epoch - 273ms/step

Epoch 6/12

326/326 - 88s - loss: 0.1425 - accuracy: 0.9484 - val_loss: 0.3055 - val_accuracy: 0.9375 - 88s/epoch - 269ms/step

Epoch 7/12

326/326 - 88s - loss: 0.1343 - accuracy: 0.9473 - val_loss: 0.3967 - val_accuracy: 0.8125 - 88s/epoch - 269ms/step

Epoch 8/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets
326/326 - 91s - loss: 0.1349 - accuracy: 0.9479 - val_loss: 0.2722 - val_accuracy: 0.8125 - 91s/epoch - 280ms/step

Epoch 9/12

326/326 - 89s - loss: 0.1247 - accuracy: 0.9503 - val_loss: 0.3647 - val_accuracy: 0.8125 - 89s/epoch - 274ms/step

Epoch 10/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets
326/326 - 91s - loss: 0.1099 - accuracy: 0.9551 - val_loss: 0.2673 - val_accuracy: 0.8750 - 91s/epoch - 280ms/step

Epoch 11/12

INFO:tensorflow:Assets written to: ./saved_models/bestmodel_0/assets
326/326 - 90s - loss: 0.1291 - accuracy: 0.9519 - val_loss: 0.2562 - val_accuracy: 0.9375 - 90s/epoch - 276ms/step

Epoch 12/12

326/326 - 90s - loss: 0.1182 - accuracy: 0.9549 - val_loss: 0.3900 - val_accuracy: 0.8750 - 90s/epoch - 277ms/step

```
In [48]: model_3 = tf.keras.models.load_model('./saved_models/bestmodel_'+str(count))
model_3.evaluate(ds_test)
```

```
39/39 [=====] - 9s 217ms/step - loss: 0.4252 - accur
acy: 0.8670
```

```
Out[48]: [0.4252314567565918, 0.8669871687889099]
```

```
In [50]: #Confusion Matrix and Classification Report
Y_pred_3 = model_3.predict_generator(ds_test, 40)
y_pred_3 = np.argmax(Y_pred_3, axis=1)
print('Confusion Matrix')
print(confusion_matrix(ds_test.classes, y_pred_3))
print('Classification Report')
target_names = ['Normal', 'Pneumonia']
print(classification_report(ds_test.classes, y_pred_3, target_names=target_names))
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: UserWarning:
`Model.predict_generator` is deprecated and will be removed in a future versi
on. Please use `Model.predict`, which supports generators.
```

```
WARNING:tensorflow:Your input ran out of data; interrupting training. Make su
re that your dataset or generator can generate at least `steps_per_epoch * ep
ochs` batches (in this case, 40 batches). You may need to use the repeat() fu
nction when building your dataset.
```

```
Confusion Matrix
```

```
[[234  0]
 [390  0]]
```

```
Classification Report
```

	precision	recall	f1-score	support
Normal	0.38	1.00	0.55	234
Pneumonia	0.00	0.00	0.00	390
accuracy			0.38	624
macro avg	0.19	0.50	0.27	624
weighted avg	0.14	0.38	0.20	624

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131
8: UndefinedMetricWarning: Precision and F-score are ill-defined and being se
t to 0.0 in labels with no predicted samples. Use `zero_division` parameter t
o control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131
8: UndefinedMetricWarning: Precision and F-score are ill-defined and being se
t to 0.0 in labels with no predicted samples. Use `zero_division` parameter t
o control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:131
8: UndefinedMetricWarning: Precision and F-score are ill-defined and being se
t to 0.0 in labels with no predicted samples. Use `zero_division` parameter t
o control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

Inferences and Comments

Explanations

- The results of the Adam optimizer are generally better than every other optimization algorithms, have faster computation time, and require fewer parameters for tuning. Because of all that, Adam is recommended as the default optimizer for most of the applications.
- The main model that used Adam optimizer and 4 hidden layers gave ~90% accuracy whereas the other 2 models gave a lower accuracy. The model that used SGD and 4 hidden layers gave ~82% accuracy. The model that used Adam optimizer and 3 hidden layers gave ~87% accuracy.